

Low-rank Updates in Statistical Physics Applications

**Phani Kumar Nukala
Oak Ridge National Laboratory**

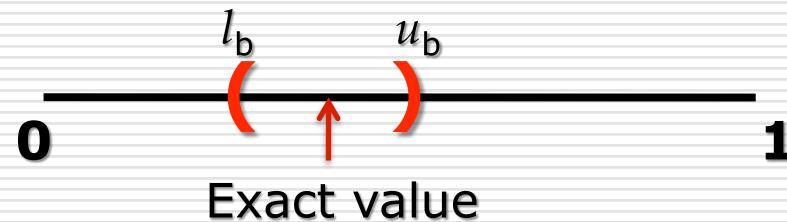
JSTAT P03029 (2010)
Phys. Rev. B 80, 195111 (2009)
J. Chem. Phys. 130(20), 204105 (2009)
JSTAT P08007 (2007)
J. Phys. A 37, 2093 (2004)
J. Phys. A 36, 11403 (2003)

Motivation

- In a typical MC simulation, computing the transition probability is the most computationally intensive step.
- In many MC models, transition probability is dependent on
 - Computing eigenvalues of Hamiltonian matrix
 - Determinant of Green's function matrix
 - or some basic linear algebra problem
- Naïve repetitive computation is a challenge even for modern supercomputers since a large number of MC steps are required to solve the problem.
 - Linear algebra sub-problem does not scale well
 - Computational complexity increases as $O(N^3)$
- **An important feature of these problems is that successive matrices differ by a low-rank perturbation**

Motivation

- Since these matrices differ by a low-rank,
 - **Can we devise algorithms that use information from previous step to EXACTLY compute the transition probability of current step?**
 - **Can we find tight bounds for this probability so that we can often avoid exact computation of transition probability?**



$$\text{Speedup} = 1/(u_b - l_b) - 1$$

Problem Definition

- Given $\det(f(\mathbf{A}_k))$, find $R_k = \frac{\det(f(\mathbf{A}_{k+1}))}{\det(f(\mathbf{A}_k))} \quad \forall k$

where $\mathbf{A}_{k+1} = \mathbf{A}_k + \alpha_k \mathbf{u}_k \mathbf{v}_k^T$

- Model problems
 - Spin-fermion systems

$$f(\mathbf{A}_k) = \mathbf{I} + e^{\beta \mathbf{A}_k}$$

- Strongly correlated systems

$$f(\mathbf{A}_k) = \mathbf{A}_k$$

- Quantum Monte Carlo (QMC)

$$f(\mathbf{A}_k) = \mathbf{A}_k$$

- Related Problems

- Discrete fracture (sparse & iterative)

$$\mathbf{A}_k \mathbf{x}_k = \mathbf{b}_k$$

Outline

- High-temperature superconductivity model (HTSC)
 - Delayed update for Hubbard Model
 - Improved sub-matrix algorithm
- Colossal Magneto-Resistance model (CMR)
 - Low-rank update algorithm for Spin-Fermion Model
 - Bounds for transition probabilities
- Statistical physics of fracture model
 - Recycling Krylov CG
 - Sparse direct solvers

Application: Hubbard Model for High-Temperature Superconductivity

Phys. Rev. B 80, 195111 (2009)
J. Chem. Phys. 130(20), 204105 (2009)

Hubbard Model

$$H = -t \sum_{\langle i,j \rangle} c_{i\sigma}^\# c_{j\sigma} + U \sum_i n_{i\uparrow} n_{i\downarrow}$$

- Transform interaction term with auxiliary HHS spin fields
- These spin fields are integrated using MC
- Local MC move is accepted based on

$$r' = \frac{\rho(\psi')}{\rho(\psi)} = \frac{\det \mathbf{G}_\uparrow^{-1}(\psi') \det \mathbf{G}_\downarrow^{-1}(\psi')}{\det \mathbf{G}_\uparrow^{-1}(\psi) \det \mathbf{G}_\downarrow^{-1}(\psi)} = R_\uparrow R_\downarrow,$$

where

$$R_\sigma = \frac{\det \mathbf{G}_\sigma^{-1}(\psi')}{\det \mathbf{G}_\sigma^{-1}(\psi)}.$$

Hubbard Model

- Repetitive computation of determinant when the Green's function matrix undergoes rank-1 updating
- Given $\det(\mathbf{G}_k)$, find $R_k = \frac{\det(\mathbf{G}_{k+1})}{\det(\mathbf{G}_k)}$

where $\mathbf{G}_{k+1} = \mathbf{G}_k + \alpha_k \mathbf{u}_k \mathbf{v}_k^T$

$$\mathbf{u}_k = \alpha_k [\mathbf{G}_k(:, p) - \mathbf{e}_p]$$

$$\mathbf{v}_k = \mathbf{G}_k(p, :)$$

$$\alpha_k = \frac{\gamma_k}{[1 + \gamma_k (1 - \mathbf{G}_k(p, p))]}$$

$$R_k = \frac{1}{[1 + \gamma_k (1 - \mathbf{G}_k(p, p))]}$$

Updating G over each step poses a significant computational problem

$O(N^2)$

Algorithm 1: Delayed Update

- Given \mathbf{G}_0 , set $\mathbf{d}_0 = \text{diag}(\mathbf{G}_0)$
- Initialize $\mathbf{u}_k = \mathbf{G}_0(:, p)$ and $\mathbf{v}_k = \mathbf{G}_0(p, :)$

- Compute

$$R_k = \frac{1}{[1 + \gamma_k(1 - \mathbf{d}_k(p))]}$$

- Update

f or i = 0 : k - 1,

$$\mathbf{u}_k = \mathbf{u}_k + \mathbf{v}_i(p) \mathbf{u}_i$$

$$\mathbf{v}_k = \mathbf{v}_k + \mathbf{u}_i(p) \mathbf{v}_i$$

end

- Update diagonal

$$\mathbf{u}_k = \alpha_k (\mathbf{u}_k - \mathbf{e}_p)$$

□ 1.35PFlops/s on
Cray Jaguar at
ORNL.

$$\mathbf{d}_k = \mathbf{d}_k + \mathbf{u}_k \circ \mathbf{v}_k$$

□ 2008 Gordon Bell
Award

$O(kN)$

Algorithm 2: Improved Method

- Instead of updating G , start working with $A = G^{-1}$

$$G_{k+1} = G_k + \alpha_k \mathbf{u}_k \mathbf{v}_k^T \quad \text{with} \quad \mathbf{u}_k = \alpha_k [G_k(:, p) - \mathbf{e}_p]$$



$$\mathbf{v}_k = G_k(p, :)$$

$$A_{k+1} = A_k + \gamma_k [A_k(:, p) - \mathbf{e}_p] \otimes \mathbf{e}_p$$

$$A_{k+1} = A_k + \begin{bmatrix} \gamma_k \\ \times \\ \times \\ \times \\ \times \end{bmatrix} - \begin{bmatrix} p \\ -\gamma_k \\ - \\ - \\ - \end{bmatrix}$$

□ For $k+1$ steps,

$$\mathbf{A}_{k+1} = \mathbf{A}_0 + \sum_{j=0}^k \gamma_j [\mathbf{A}_0(:, \mathbf{p}(j)) - \mathbf{e}_{\mathbf{p}(j)}] \otimes \mathbf{e}_{\mathbf{p}(j)}$$

$$= \tilde{\mathbf{A}}_k - \sum_{j=0}^k \gamma_j \mathbf{e}_{\mathbf{p}(j)} \otimes \mathbf{e}_{\mathbf{p}(j)} = \tilde{\mathbf{A}}_k - \mathbf{U}\mathbf{V}^T$$

$$\mathbf{A}_{k+1} = \mathbf{A}_0 + \left[\begin{array}{cc} \gamma_0 & \times \\ \times & \gamma_k \\ \times & \times \\ \times & \dots & \times \\ \times & \times \\ \times & \times \end{array} \right] - \left[\begin{array}{cc} \mathbf{p}(0) & \mathbf{p}(k) \\ | & | \\ - & \gamma_0 & - & - & - \\ | & | & | & | \\ - & - & - & \gamma_k \\ | & & | & | \end{array} \right]$$

$$\det(\mathbf{A}_{k+1}) = \det(\tilde{\mathbf{A}}_k) \det(\mathbf{I} - \mathbf{V}^T \tilde{\mathbf{G}}_k \mathbf{U})$$

$$\tilde{\mathbf{A}}_{k+1} = \mathbf{A}_0 + \begin{bmatrix} \gamma_0 & & \gamma_k \\ \times & & \times \\ \times & & \times \\ \times & \dots & \times \\ \times & & \times \\ \times & & \times \end{bmatrix} \rightarrow \det(\tilde{\mathbf{A}}_k) = \left[\prod_{j=0}^k (1 + \gamma_j) \right] \det(\mathbf{A}_0)$$

$$\det(\mathbf{I} - \mathbf{V}^T \tilde{\mathbf{G}}_k \mathbf{U}) = (-1)^{k+1} \left[\prod_{j=0}^k \frac{\gamma_j}{1 + \gamma_j} \right] \det(\Gamma_k)$$

where

$$\Gamma_k = \mathbf{G}_0(\mathbf{p}) - \begin{bmatrix} (1 + \gamma_0) & & & \\ & \ddots & & \\ & & (1 + \gamma_k) & \\ & & & \ddots \end{bmatrix}$$

$$\mathbf{G}_0(\mathbf{p}) = \begin{bmatrix} \mathbf{G}_0(\mathbf{p}(0), \mathbf{p}(0)) & \cdots & \mathbf{G}_0(\mathbf{p}(0), \mathbf{p}(k)) \\ \vdots & \ddots & \vdots \\ \mathbf{G}_0(\mathbf{p}(k), \mathbf{p}(0)) & \cdots & \mathbf{G}_0(\mathbf{p}(k), \mathbf{p}(k)) \end{bmatrix}$$

□ Symbolically,

$$\boldsymbol{\Gamma}_k = \begin{bmatrix} \boldsymbol{\Gamma}_{k-1} & \mathbf{s} \\ \mathbf{w}^T & d \end{bmatrix} = \begin{bmatrix} \mathbf{L}_{k-1} & 0 \\ \mathbf{x}^T & 1 \end{bmatrix} \begin{bmatrix} \mathbf{U}_{k-1} & \mathbf{y} \\ 0 & \beta \end{bmatrix}$$

where

$$\mathbf{s} = \mathbf{G}_0(\mathbf{p}(0) : \mathbf{p}(k-1), \mathbf{p}(k))$$

$$\mathbf{w}^T = \mathbf{G}_0(\mathbf{p}(k), \mathbf{p}(0) : \mathbf{p}(k-1))$$

$$d = \mathbf{G}_0(\mathbf{p}(k), \mathbf{p}(k)) - \frac{1 + \gamma_k}{\gamma_k}$$

$$\det(\boldsymbol{\Gamma}_k) = \beta \det(\boldsymbol{\Gamma}_{k-1})$$

where

$$\beta = d - \mathbf{x}^T \mathbf{y}$$

$$\mathbf{L}_{k-1} \mathbf{y} = \mathbf{s}$$

$$\mathbf{U}_{k-1}^T \mathbf{x} = \mathbf{w}$$



$$O(k^2)$$

Summarizing ... $O(k^2)$ algorithm

$$\det(\mathbf{A}_{k+1}) = \det(\tilde{\mathbf{A}}_k) \det(\mathbf{I} - \mathbf{V}^T \tilde{\mathbf{G}}_k \mathbf{U})$$

$$\det(\tilde{\mathbf{A}}_k) = \left[\prod_{j=0}^k (1 + \gamma_j) \right] \det(\mathbf{A}_0)$$

$$\det(\mathbf{I} - \mathbf{V}^T \tilde{\mathbf{G}}_k \mathbf{U}) = (-1)^{k+1} \left[\prod_{j=0}^k \frac{\gamma_j}{1 + \gamma_j} \right] \det(\Gamma_k)$$

$$\det(\Gamma_k) = \beta \det(\Gamma_{k-1})$$

$$\rightarrow \boxed{\det(\mathbf{A}_{k+1}) = -\beta \gamma_k \det(\mathbf{A}_k)} \rightarrow R_k = -\frac{1}{\beta \gamma_k}$$

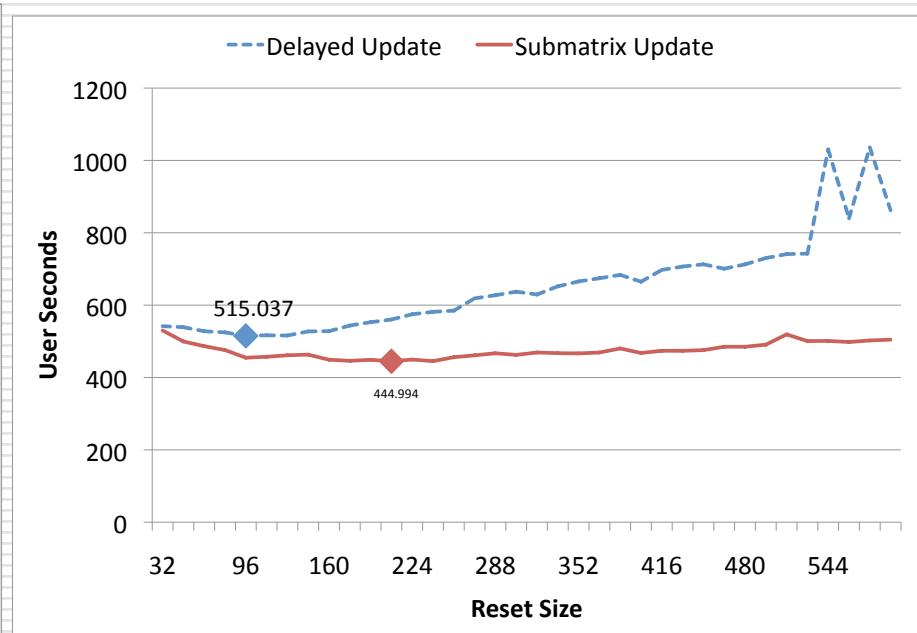
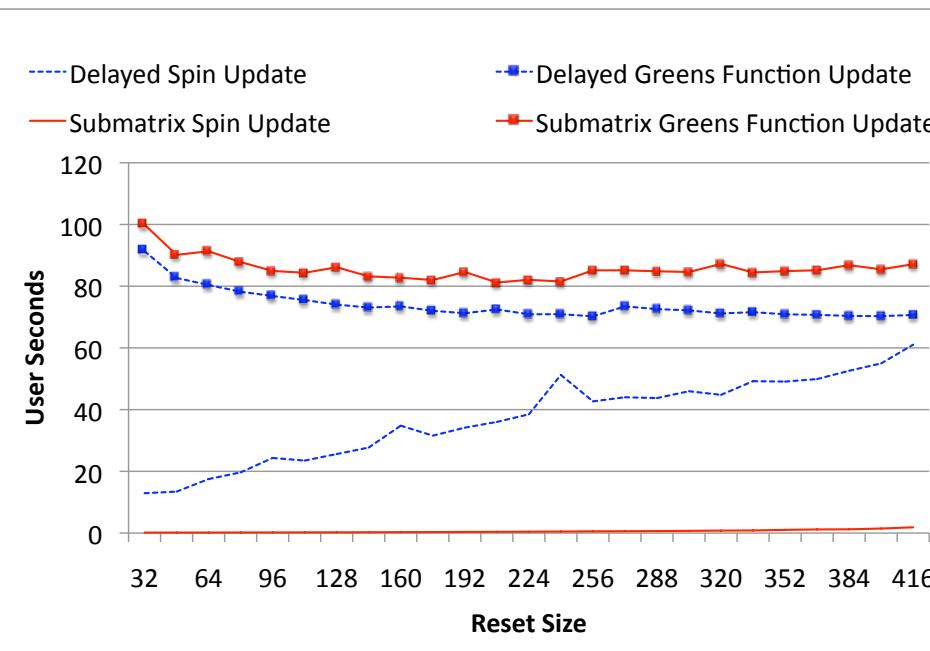
Numerical Results: Hubbard Model

- Given $\det(\mathbf{G}_k)$, find $R_k = \frac{\det(\mathbf{G}_{k+1})}{\det(\mathbf{G}_k)}$, such that
 - $\mathbf{G}_{k+1} = \mathbf{G}_k + \alpha_k \mathbf{u}_k \mathbf{v}_k^T$
 - $\mathbf{u}_k = \alpha_k [\mathbf{G}_k(:, p) - \mathbf{e}_p]$
 - $\mathbf{v}_k = \mathbf{G}_k(p, :)$
- Initialized with a random matrix
- Updates, $m = N/10$

N	Full	Delayed	Recursive	New
1000	7.02	0.28	0.09	0.015
3000	195	9.86	2.31	0.49

Numerical Results: Hubbard Model

- 16 site dynamic cluster QMC
- 2D Hubbard model
- Hopping t , and $U = 4t$; inverse temperature = $40/t$



Application: Colossal Magneto-resistance and Spin-Fermion Systems

Spin-Fermion Systems: Double Exchange Model

$$H(\psi) = c^\# A(\psi) c$$

$$\begin{aligned} \mathbf{A}_{i+\sigma M, j+\sigma' M}^\alpha &= t_{ij} \delta_{\sigma, \sigma'} + (1 - 2\sigma) J_H \cos(\theta_i^\alpha) \delta_{i,j} \delta_{\sigma, \sigma'} \\ &\quad + J_H \exp(i\phi_i^\alpha (1 - 2\sigma)) \sin(\theta_i^\alpha) \delta_{i,j} \delta_{\sigma, 1-\sigma'}. \end{aligned}$$

- Boltzmann weight $P(\psi)$ for field configuration ψ is given by

$$\begin{aligned} P(\psi) &= \exp(-S_{\text{eff}}(\psi))/Z \\ Z &= \sum_{\psi} \exp(-S_{\text{eff}}(\psi)) \\ S_{\text{eff}} &= \sum_{\nu} F(\lambda_{\nu}(\psi)) \\ F(\lambda) &= -\log(1 + \exp(-\beta(\lambda - \mu))) \end{aligned}$$

Spin-Fermion Systems

□ Simulation of colossal magnetoresistance using spin-fermion systems poses the following problem:

□ Given $\det(\mathbf{I} + e^{\beta \mathbf{A}_k})$, find $\det(\mathbf{I} + e^{\beta \mathbf{A}_{k+1}})$

where $\mathbf{A}_{k+1} = \mathbf{A}_k + \rho \mathbf{u} \mathbf{u}^T$ for all $k = 0, 1, 2, \dots$

■ Updating all the eigenvalues of \mathbf{A}_{k+1} based on the eigenvalues of \mathbf{A}_k

Spin-fermion Systems

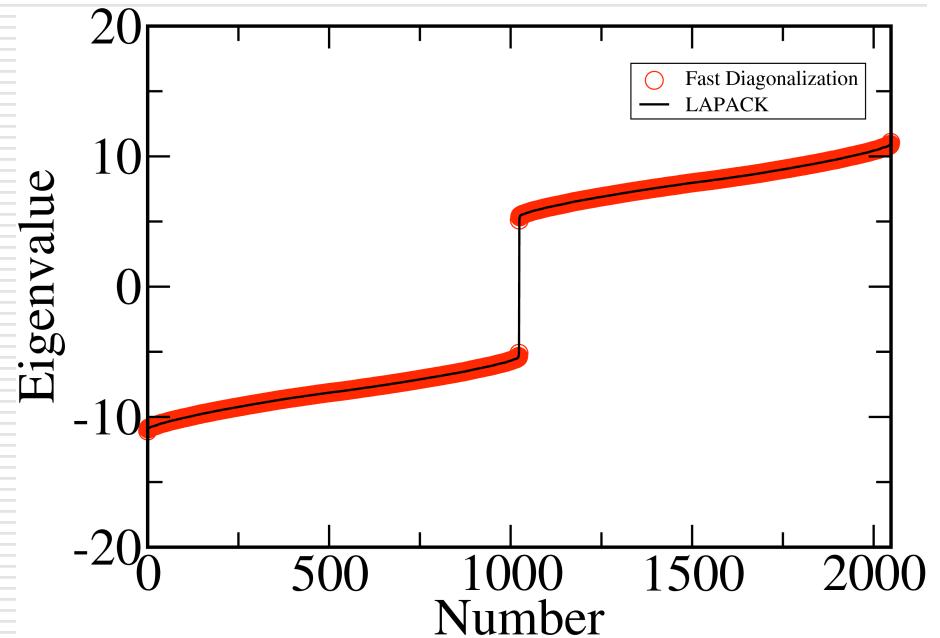
Comparison of CPU times taken for 10 steps

Size (N)	LAPACK (sec)	Low-rank update (sec)
288	0.6	0.34
800	16.1	2.5
1152	60.1	9.7
2048	298.5	32.1

$O(N^3)$

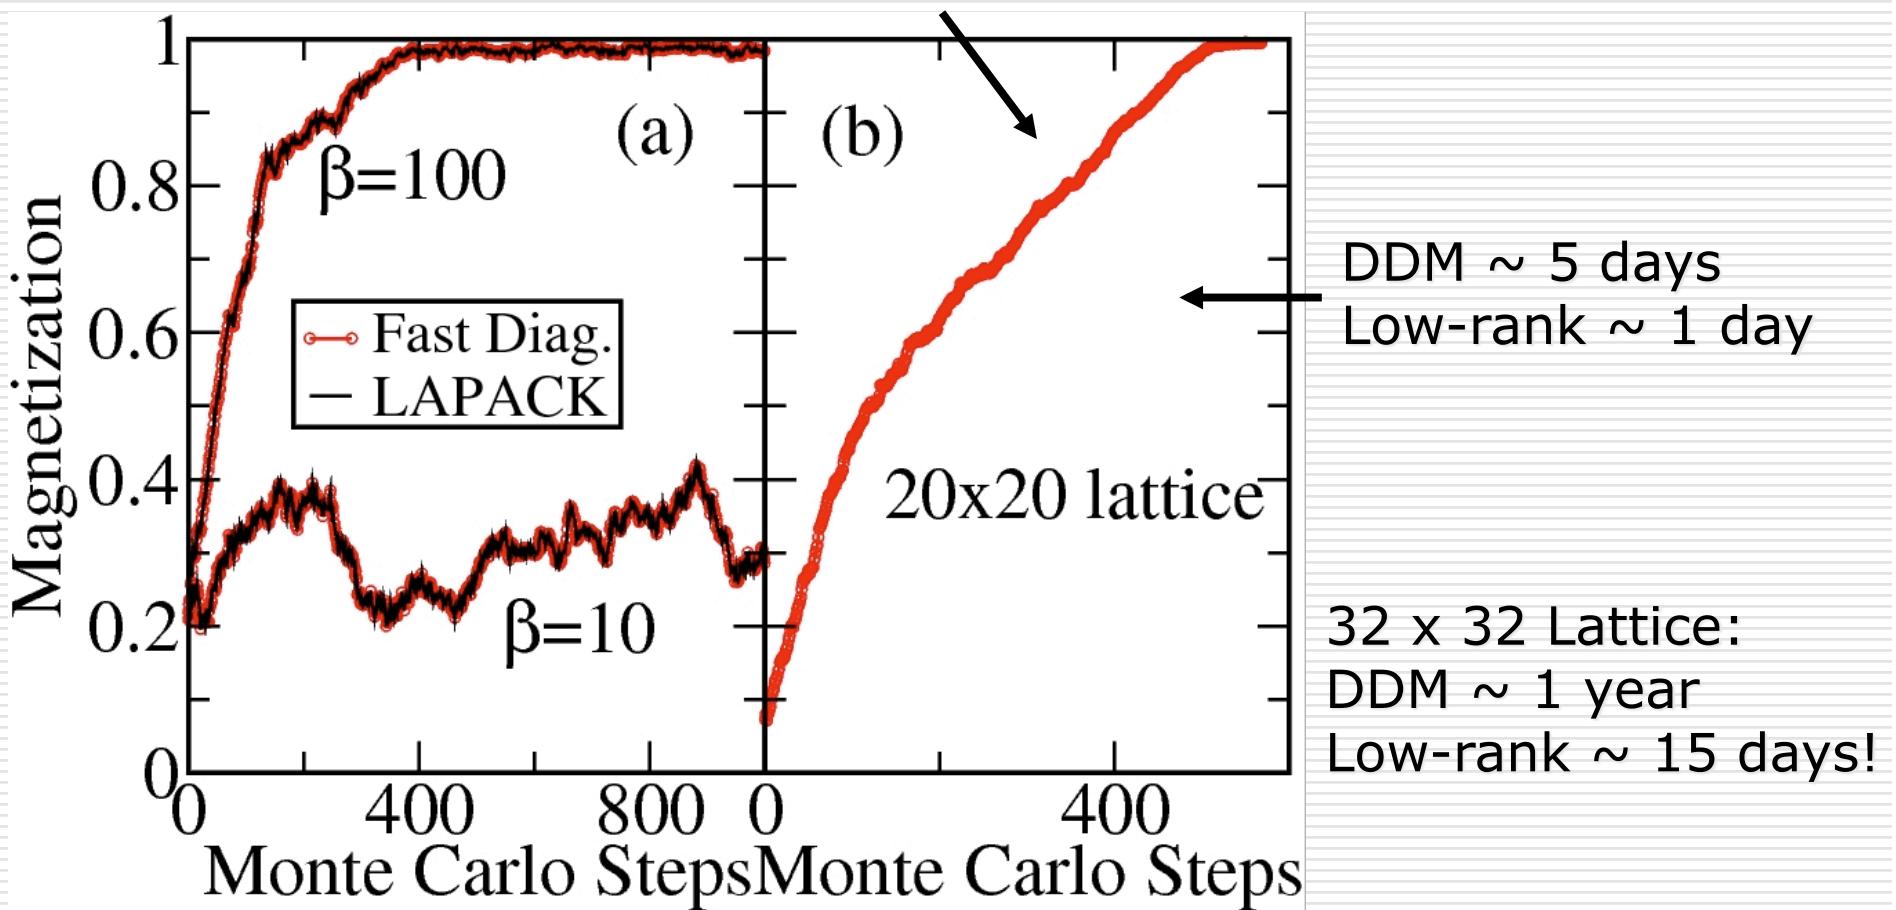
$O(N^2)$

- Fast Multipole Method further reduces the computational complexity to $O(N \log N)$
- Excellent accuracy of eigen spectrum even after many updates



Large System Simulation

- Simulation not readily accessible to traditional method of direct diagonalization (DDM) during each step



Estimating Bounds for Transition Probabilities in QMC

- Given $\det(f(\mathbf{A}_k))$, find bounds C_l and C_u such that

$$C_l \leq R_k = \frac{\det(f(\mathbf{A}_{k+1}))}{\det(f(\mathbf{A}_k))} \leq C_u$$

where $\mathbf{A}_{k+1} = \mathbf{A}_k + \alpha_k \mathbf{u}_k \mathbf{v}_k^T$

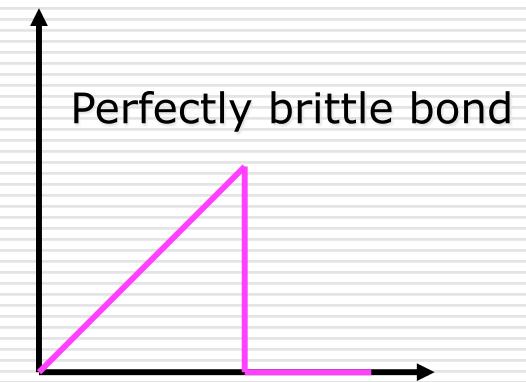
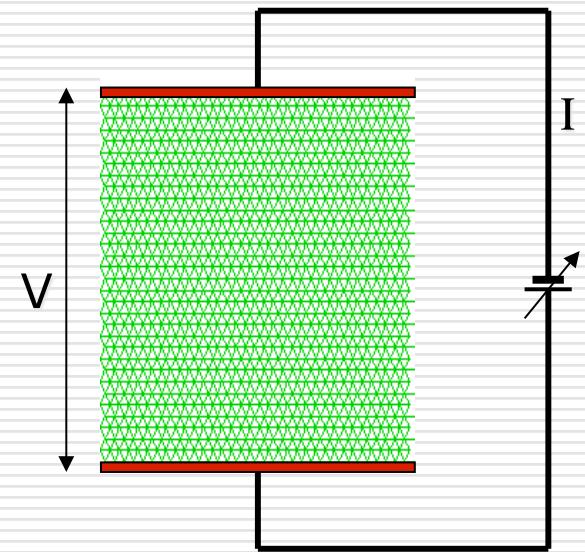
- Accurate bounds would significantly speedup Monte Carlo simulation

Application: Statistical Physics of Fracture

JSTAT P03029 (2010)
J. Phys. A 37, 2093 (2004)
J. Phys. A 36, 11403 (2003)

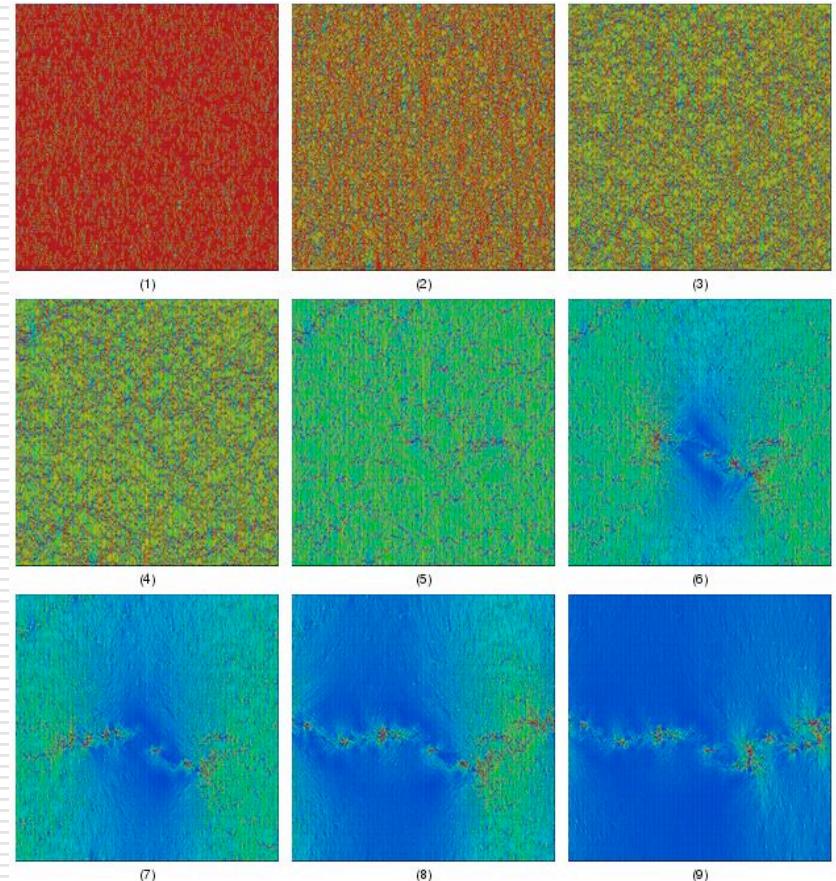
Random Thresholds Models

- For each bond, assign unit conductance and the thresholds are prescribed based on a random thresholds distribution
- The bond breaks irreversibly whenever the current (stress) in the fuse exceeds the prescribed thresholds value
- Currents (stresses) are redistributed instantaneously
- The process of breaking one bond at a time is repeated until the lattice falls apart



Fracture of a 2D lattice system

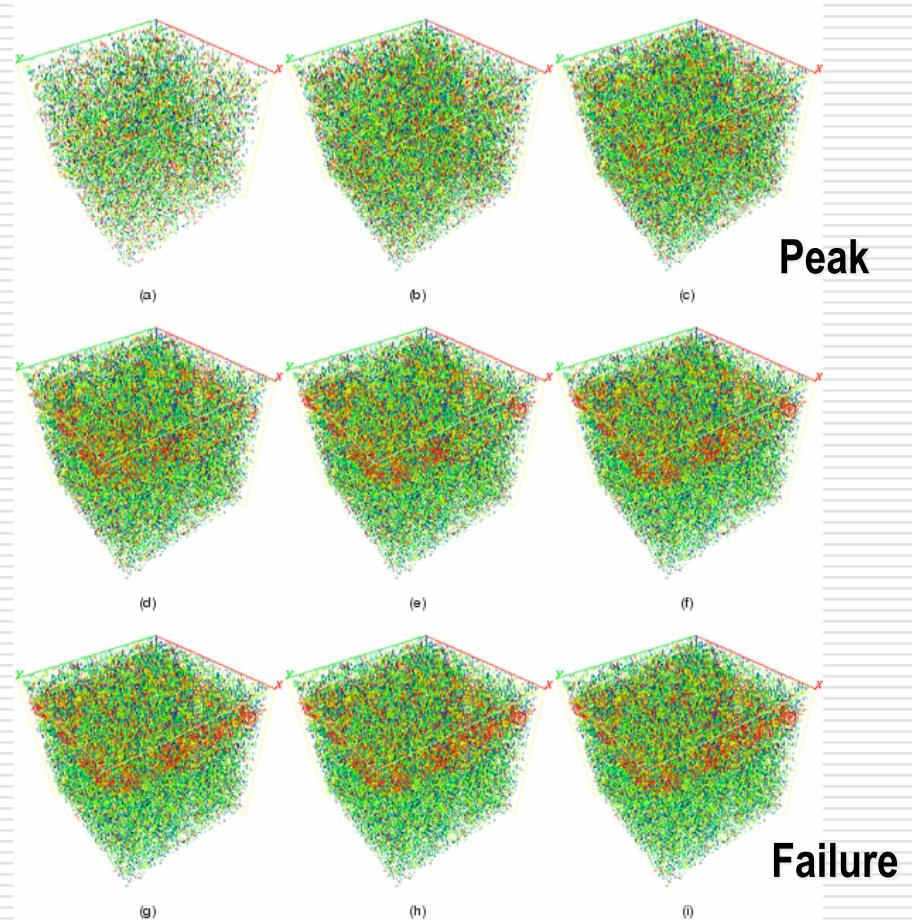
- CPU $\sim O(L^{4.5})$
- Capability issue: Previous simulations have been limited to a system size of $L = 128$
- Largest 2D lattice system ($L = 1024$) analyzed for investigating fracture and damage evolution
- Effective computational gain ~ 80 times



Stress redistribution in the lattice due to progressive damage/crack propagation

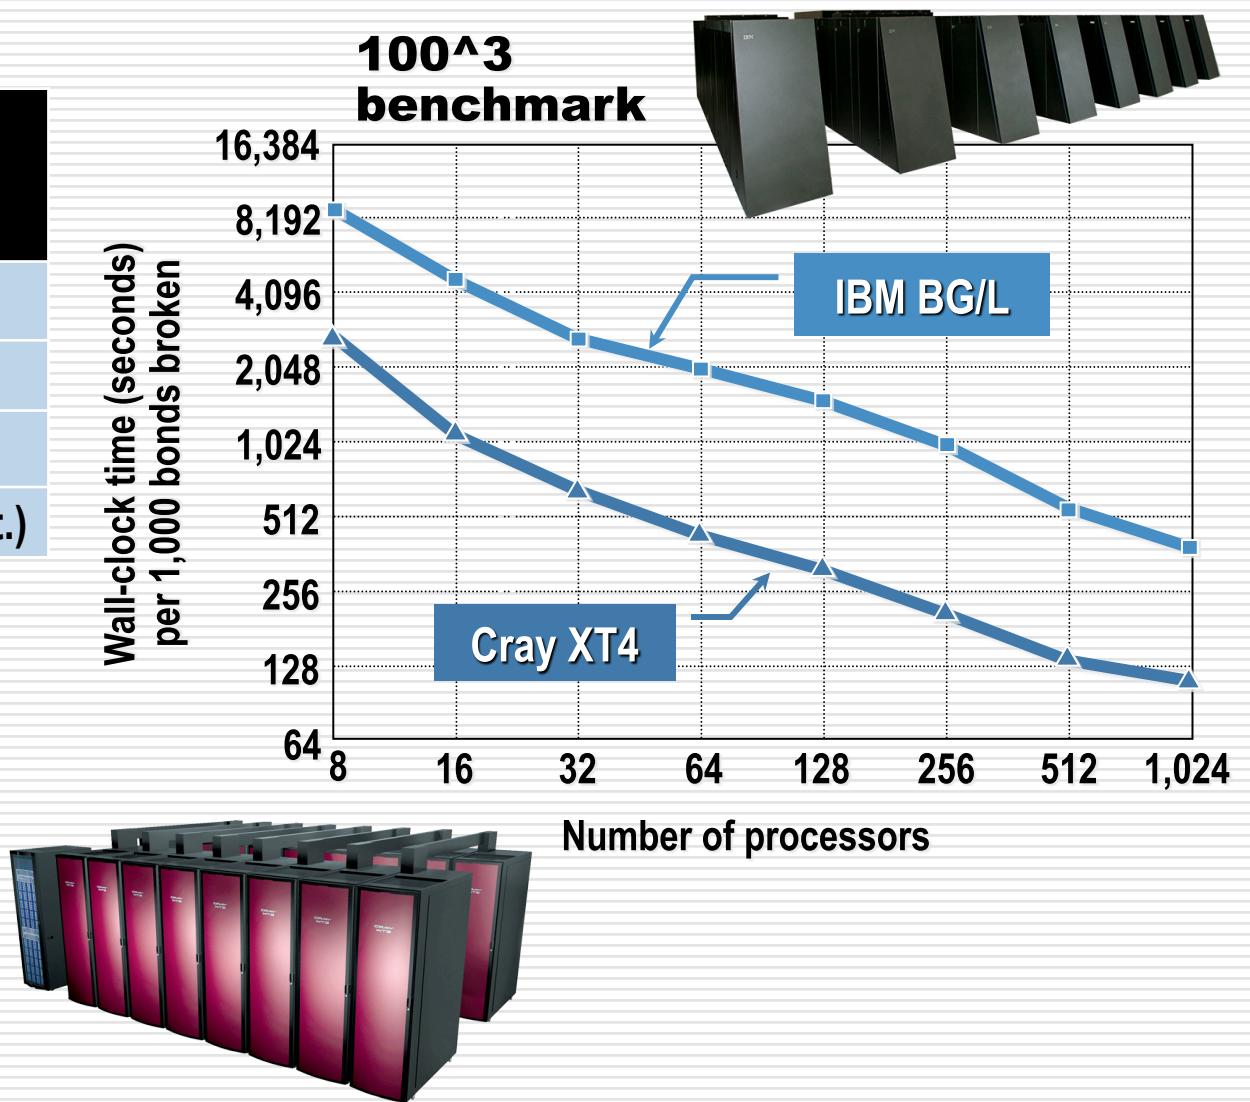
Fracture of 3D lattice system

- CPU $\sim O(L^{6.5})$
- Largest cubic lattice system analyzed for investigating fracture and damage evolution in 3D ($L = 64$)
- On a single processor, a 3D system of size $L = 64$ requires 15 days of CPU time!



High-performance computing

High-performance computing	Processing time
L = 64 on 128	3 hours
L = 100 on 1024	12 hours
L = 128 on 1024	3 days
L = 200 on 2048	20 days (est.)



Recycling Conjugate Gradients for Fracture Modeling

- Recycling approximate invariant subspaces (lowest eigenvectors) between linear systems
- Projection step in CG to enforce all estimates are orthogonal to $\text{span}(u_1, u_2, \dots, u_k)$
- Algorithm to generate new recycling space for next system by using conjugate directions to update estimates of harmonic Ritz vectors
- Development of model specific optimization in low rank update ($A \leftarrow A + \sigma vv'$)
- High cost for accurate computation of invariant subspace can be amortized over many solves

Summary

- For Monte Carlo applications,
 - a low-rank updating scheme combined with bounds on transition probabilities can significantly speed up computation.
 - Typical speedups $\sim 5\text{-}50$ times
- Applications
 - Hubbard model for strongly correlated systems (QMC)
 - Low-rank schemes and bounds on transition probability
 - Spin-Fermion models
 - Repetitive eigen-decompositions and bounds on transition probability
 - Discrete Lattice models of fracture
 - Recycled CG
 - Sparse direct solvers

Monte Carlo Simulation Scheme

- Estimate thermodynamic properties
- Start with an initial configuration (chosen at random)
- During each MC step,
 - Propose a local change
 - **Determine the probability to change**
 - **Either a Boltzmann weight or some other positive unit normalized measure**
 - Accept or reject change (Metropolis or heat-bath)
- Estimate time averages during measurement stage, and by ergodic theory, equate them to thermodynamics properties